

Human Posture Analysis and Correction

Krishna Patil ¹, Prajwal Billade ², Samraddhi Rathor ³, Maitreyee Joshi ⁴, Prof. P.T. Kohok ⁵, Dr. G.V. Kale ⁶

¹Student; SCTR's Pune Institute of Computer Technology, (Computer Engineering), Pune, Maharashtra, India, krishnap28113@gmail.com

²Student; SCTR's Pune Institute of Computer Technology, (Computer Engineering), Pune, Maharashtra, India, prajwalbillade@gmail.com

³Student; SCTR's Pune Institute of Computer Technology, (Computer Engineering), Pune, Maharashtra, India, rathorsamraddhi@gmail.com

⁴Student; SCTR's Pune Institute of Computer Technology, (Computer Engineering), Pune, Maharashtra, India, maitreyeejoshi@gmail.com

⁵Assistant Professor; SCTR's Pune Institute of Computer Technology, (Computer Engineering), Pune, Maharashtra, India, ptkohok@pict.edu

⁶Associate Professor; SCTR's Pune Institute of Computer Technology, (Computer Engineering), Pune, Maharashtra, India, gvkale@pict.edu

Abstract

The analysis and correction of human posture have attracted considerable research interest in recent years, primarily due to their importance in preventing musculoskeletal disorders (MSDs) and improving overall well-being. This study presents a comprehensive review of over 80 contemporary works in the domains of computer vision and machine learning, with particular emphasis on hybrid approaches, sensor-assisted systems, and advanced 2D/3D pose estimation techniques, including frameworks like MediaPipe and OpenPose. We also examine widely used benchmark datasets such as Yoga-82 and MPED, alongside a custom-built dataset comprising 500 video samples. These videos, captured using standard consumer-grade devices, encompass a wide variety of body types and dynamic movements across five exercise categories, including activities like squats and deadlifts. Through the development and implementation of a posture classification model, we achieved an average accuracy of 82.5%, showcasing the practicality of performing real-time inference on devices with constrained computational resources, such as everyday laptops. However, the study also highlights persistent challenges, notably in dealing with occlusions, meeting real-time processing demands, and addressing privacy concerns related to sensitive user data. In exploring potential solutions, emerging trends like federated learning and synthetic data generation, which offer promising avenues for improving scalability and mitigating ethical issues were discussed. Finally, providing actionable recommendations to advance posture analysis applications in fields such as healthcare, workplace ergonomics, and AI-assisted fitness was ensured. Our findings underscore the need for greater standardization, user-centered design considerations, and scalable deployment strategies to foster wider adoption of posture monitoring technologies.

Keywords: Anthropometry, Scene Understanding, Activity Recognition, Scene Anomaly Detection, Visual Inspection

1. Introduction

Maintaining correct posture is crucial for overall health and well-being, yet prolonged periods of sitting can negatively impact body alignment. Habitual poor posture—such as slouching or improper seating ergonomics—can lead to musculoskeletal disorders (MSDs), impairing daily functioning and reducing quality of life. Traditionally, posture evaluation relied on manual clinical assessments and physical examinations. However, advancements in computer vision and machine learning have made it possible to automate posture analysis with high accuracy and in real-time.

The integration of artificial intelligence (AI) and deep learning has significantly advanced the field of human posture estimation. Modern techniques primarily use vision-based approaches to detect anatomical keypoints and

assess postural alignment. Early pose estimation methods depended on handcrafted features like Histogram of Oriented Gradients (HOG) and Scale-Invariant Feature Transform (SIFT) for landmark detection. While these methods offered initial solutions, they struggled with generalization in uncontrolled, real-world environments. The advent of deep learning, particularly through convolutional neural networks (CNNs), has dramatically improved keypoint detection accuracy.

Contemporary models such as OpenPose, DeepPose, and MediaPipe now enable efficient real-time tracking of human posture across diverse applications, including healthcare monitoring, fitness coaching, ergonomic assessments, and physical rehabilitation. Despite these technological advancements, significant challenges remain. Issues such as limited availability of diverse datasets, real-time multi-person tracking difficulties, and inaccuracies caused by occlusions continue to hinder system robustness in everyday scenarios. While current solutions perform well under controlled laboratory settings, their reliability in naturalistic environments still requires improvement. Addressing these challenges is essential for making posture estimation and correction technologies more practical and widely applicable. This work aims to provide a comprehensive overview of human pose estimation methods, analyze their current applications, highlight persistent challenges, and propose future research directions to enhance their effectiveness and scalability.

2. Methodology used in this research work

2.1 Vision-Based Posture Analysis

The emerging trends in posture analysis systems are increasingly leveraging computer vision and artificial intelligence (AI) to enhance the accurate detection of postural defects. Among the key features, real-time feedback plays a crucial role by providing instant results for all body positions, allowing for the immediate detection of adverse postures [9]. Additionally, these systems offer automated assessments, which significantly improve clinical management and facilitate more effective fitness monitoring [15]. By automating posture evaluation, these technologies not only reduce the need for manual intervention but also ensure consistent and objective analysis, thus contributing to better health outcomes and preventive care. The following are the role of key models in posture analysis:

2.1.1 OpenPose

A foundation in multi-person pose estimation, OpenPose has been widely used in rehabilitation and group fitness applications [12]. However, real-time deployment in resource-constrained environments such as mobile apps is limited by its dependency on GPU acceleration. In dynamic exercises like squats, occlusion robustness can be improved by combining OpenPose with temporal models like LSTM [19].

2.1.2 MediaPipe

MediaPipe's lightweight architecture enables real-time posture correction in consumer-facing apps (e.g., yoga and weightlifting trainers), which were designed for edge devices. Ko et al. [1] showed its utility in powerlifting by combining it with YOLOv5 for person detection. The lack of depth perception is still a barrier for 3D biomechanical analysis.

2.1.3 DeepPose

As one of the first deep learning models for pose estimation, DeepPose served as a foundation for later architectures like Transformers and ResNet [12]. Despite being computationally intensive, its foundational approach inspired hybrid models that combine CNNs with geometric constraints to estimate joint angles [6].

2.1.4 ResNet-101

It is ideal for clinical applications that require precision, such as spinal deformity screening, due to model's high accuracy [2]. However, its use in domains like powerlifting is limited by its dependence on large datasets (e.g., Yoga-82), where labeled data is scarce.

Table 1. Comparative analysis of existing pose estimation models

Model	Accuracy	Real-Time	Limitations
OpenPose	High	Yes	Occlusion issues
MediaPipe	High	Yes	Limited depth estimation
DeepPose	Moderate	No	High computational cost
ResNet-101	Very High	No	Require large datasets

2.2 Posture recognition methods

2.2.1 Sensor-Based Approaches

Sensor-based posture recognition tracks a person's movement and body position using wearable technologies such as gyroscopes and accelerometers. These techniques have been widely adopted in fields like industrial ergonomics and medical rehabilitation due to their effectiveness. One of the main advantages of sensor-based systems is their high accuracy in controlled environments, making them reliable for precise posture tracking. Additionally, they perform well in non-visual environments, such as low-light conditions, where camera-based systems might fail. However, these systems also come with certain limitations. They require physical contact with the user, which can be uncomfortable or intrusive during prolonged use. Furthermore, the cost of deploying wearable sensors on a large scale can be quite high, posing a significant challenge for widespread adoption. Kim et al. [2] developed a sensor-based spinal posture monitoring system, achieving an accuracy of 94%.

2.2.2 Vision-Based Approaches

This method employs computer vision techniques to recognize and classify postures from images or video frames. These techniques rely on models designed for posture estimation, skeletal analysis, and the detection of key body points. Vision-based recognition can be broadly categorized into two approaches. The first is 2D vision-based recognition, which extracts body landmarks from RGB images using deep learning techniques. While effective in certain scenarios, this approach faces challenges related to depth perception and occlusion, which can hinder accurate posture analysis. On the other hand, 3D vision-based recognition constructs three-dimensional body models by utilizing depth sensors such as Microsoft Kinect and LiDAR. This significantly enhances the accuracy of posture detection by providing more comprehensive spatial information. However, the use of specialized hardware makes this method more expensive and adds complexity to its deployment, posing practical limitations for widespread application. Piñero-Fuentes et al. [15] presented a 3D vision-based system for recognizing workplace posture, achieving an accuracy rate of 80%.

Table 2. Comparative analysis of existing pose estimation models

Approach	Accuracy	Hardware	Limitations
Sensor	High	Wearable Devices	Uncomfortable for users
2D Vision	Moderate	Standard RGB cameras	Occlusion issues
3D Vision	High	Depth sensors	Expensive hardware
Deep Learning	Very High	Powerful GPUs	High computational costs

2.2.3 Deep Learning-Based Approaches

Deep learning significantly facilitates the automation of posture representation learning by eliminating the need for manual feature extraction. Convolutional Neural Networks (CNNs) are commonly employed to extract relevant features from images, enabling the system to identify postural characteristics effectively. In addition, Long Short-Term Memory networks (LSTMs) are utilized to monitor variations in movement over time, which is essential for dynamic posture analysis. More recently, the incorporation of Transformers and attention mechanisms has further enhanced model performance by allowing the network to focus on important body parts while managing occlusions.

more effectively. These advancements collectively improve the accuracy and robustness of posture recognition systems, making deep learning a critical component in modern posture analysis solutions.

2.3 Dataset and Benchmarking Challenges

The availability and quality of the data sets play an essential role in the advancement of posture analysis systems. Following are the prominent datasets and their limitations, along with the role of synthetic data in addressing these challenges.

MPED:

The MultiPose Exercise Dataset (MPED), introduced by Ko et al. [1], is the first 3D dataset specifically created for powerlifting exercises such as bench press, squats, and deadlifts. This dataset captures joint coordinates during both concentric and eccentric phases of these movements by utilizing multi-angle smartphone cameras, enabling detailed biomechanical analysis. One of the key strengths of MPED is its phase-specific annotations for dynamic movements, which provide valuable insights for exercise evaluation. Additionally, by collecting data in real-world gym environments, MPED enhances the practical relevance and applicability of its findings. However, the dataset also has notable limitations. It features data from only a single experienced lifter, which restricts its ability to generalize across different body types, skill levels, and demographic groups. While multi-view camera setups help mitigate occlusion issues, challenges posed by complex gym environments, such as crowded spaces, remain unaddressed. Furthermore, the homogeneous nature of the participant demographics may lead to biased model performance for underrepresented groups, limiting its broader applicability.

Table 3. Dataset Description

Video ID	Subject ID	Exercise Type	Camera Angle	Lighting Condition	Posture Label	Mistake Type
VID_Deadlift_3	SUB_02	Deadlift	Side (Right)	Artificial Light	Incorrect	Rounded Back
VID_Squats_2	SUB_01	Squat	Front	Natural Light	Correct	None
VID_Squats_5	SUB_04	Squat	Diagonal 45°	Mixed Lighting	Incorrect	Knee Valgus
VID_Deadlift_1	SUB_03	Deadlift	Side (Left)	Natural Light	Correct	None
VID_Squats_3	SUB_02	Squat	Side (Right)	Artificial Light	Incorrect	Shallow Depth
VID_Deadlift_5	SUB_05	Deadlift	Front	Mixed Lighting	Incorrect	Excessive Forward Lean
VID_Squats_1	SUB_01	Squat	Diagonal 45°	Artificial Light	Correct	None
VID_Deadlift_2	SUB_04	Deadlift	Diagonal 45°	Natural Light	Incorrect	Rounded Back
VID_Squats_4	SUB_03	Squat	Side (Left)	Natural Light	Correct	None
VID_Deadlift_4	SUB_02	Deadlift	Front	Artificial Light	Correct	None

Yoga-82:

The Yoga-82 dataset comprises 29,000 images representing 82 different yoga postures, significantly contributing to the advancement of yoga pose identification. Its extensive coverage of static poses, such as the Tree Pose and Downward Dog, along with marked joint angles and body landmarks, enables precise classification and analysis. Despite these strengths, the dataset exhibits several limitations. It primarily focuses on static postures, neglecting transitional movements and dynamic sequences that are prevalent in real-world yoga practice. Moreover, the images were captured in controlled environments, which reduces the model's robustness when applied to settings with varying lighting conditions and complex backgrounds. Additionally, Yoga-82 centers on traditional yoga asanas, thereby overlooking modern hybrid practices like power yoga, limiting its applicability to diverse styles and contemporary fitness routines.

Synthetic Data:

Synthetic data generation tools, such as Blender and Unity, have emerged as effective solutions to address data scarcity, particularly for specialized applications like powerlifting exercises or the analysis of rare postural deformities. These tools offer significant advantages, including the ability to generate virtually infinite variations of poses, body types, and environmental conditions [12]. Moreover, they help avoid ethical concerns associated with collecting data from real users, ensuring privacy and consent-related issues are bypassed. However, there are notable limitations to synthetic data. In real-world scenarios, models trained solely on synthetic datasets often underperform due to discrepancies in lighting, textures, and joint kinematics when compared to actual human movements [12]. Furthermore, critical aspects of human motion are essential for rehabilitation applications—such as muscle fatigue and micro-tremors—are difficult to replicate accurately with synthetic avatars [6]. While benchmarks like Human3.6M [21] are commonly used to validate synthetic data, their focus on simple actions, such as walking, limits their applicability to more complex exercises and dynamic postural tasks.

3. Literature Survey

Table 4. Literature Survey

Reference	Methodology/ Technique	Dataset	Key Contribution	Limitations
Ko et al.(2024)-[1]	YOLOv5 + MediaPipe	MPED	Phase-specific feedback for power-lifting exercises.	Single-user focus, lacks body type diversity.
Kim et al.(2022)-[2]	Computer vision-based system	140 participants (Spinal posture)	94% accuracy in spinal deformity screening.	Intrusive for gym use, limited dynamic analysis.
Luangaphirom et al. (2024)-[3]	MediaPipe	Real-time weight training	Real-time counting and correction for weightlifting.	Limited to simple movements, no depth perception.
Sakshi & Saini (2024)-[4]	MediaPipe + Deep Learning	Yoga poses (Lecture Notes)	Yoga posture estimation and correction framework.	Lacks 3D analysis, small dataset.
Li et al. (2019)-[6]	Kinect V2 + BP Neural Network	Custom gym dataset	Hybrid approach combining anthropometry and deep learning.	Small dataset (n=50), limited to controlled environments.
Ding et al. (2020)-[8]	Multiple features + Rule learning	Custom dataset	Human posture recognition using feature fusion.	Moderate accuracy, not real-time.
Shan et al. (2024)-[9]	Deep Learning (CNN-LSTM)	Custom dataset	91% accuracy.	No temporal analysis, limited to 2D.
Wang & Shi (2024)-[11]	MobileNetV2	Physical education teaching	Deep learning-driven posture location in classrooms.	Limited to controlled environments, low dynamic accuracy.
Toshev & Szegedy (2014)-[12]	DeepPose (CNN)	ImageNet	Pioneered deep learning for 2D pose estimation.	High computational cost, struggles with occlusions.
Kanase et al. (2021)-[14]	OpenPose + DTW distance	Custom fitness dataset	Rule-based corrective feedback for exercises.	Pre-recorded video dependency, no real-time analysis.

Piñero-Fuentes et al. (2021)-[15]	3D CNN + Kinect	Telework-related MSDs	Posture detection system for work- place ergonomics.	Expensive hardware (Kinect), limited scalability.
Panini & Cucchiara (2003)-[16]	Machine Learning	Domotics applications	Early ML approach for posture detection in smart homes.	Outdated techniques, limited accuracy.
Chen & Yang (2020)-[19]	OpenPose + DTW	Custom fitness dataset	Real-time exercise posture correction (e.g., squats).	Limited to pre-recorded videos, lacks dynamic environments.
Pothanicker (2019)-[22]	CNN + LSTM-RNN + Attention	Custom action dataset	Human action recognition with attention mechanisms.	Limited to simple actions, small dataset.

4. System Architecture

4.1 Input Section




Fig. 1. Input Section in this research work

The input block represents the first and foundational stage of the posture monitoring and correction system. Its primary function is to acquire, process, and prepare visual data for subsequent analysis, ensuring that posture-related information is accurately captured and reliably passed through the system. This block handles the complexities of raw data acquisition and transforms it into a structured format suitable for posture estimation and classification tasks. Below is an in-depth explanation of each component of the input block and its role in the overall system architecture.

Video Input Feed (Camera / Dataset):

The entry point of the system is the video input feed, which serves as the primary source of visual data for posture analysis. This input can originate from live camera streams—such as webcams, smartphone cameras, or dedicated monitoring devices—or from pre-recorded datasets used for training, testing, and validation purposes. The quality, resolution, and frame rate of the video feed are critical factors that influence the accuracy of subsequent posture detection processes. The system is designed to handle variations in input sources, ensuring compatibility with diverse camera configurations and dataset formats.

Pre-processing (Resizing, Denoising, Normalization):

Raw video frames often contain noise, inconsistent resolutions, and lighting variations that can adversely affect the performance of downstream models. Therefore, a dedicated pre-processing module is employed to standardize and enhance the input data. The video frames are resized to a consistent resolution, ensuring uniformity in model inputs. Denoising techniques, such as Gaussian blurring or median filtering, are applied to mitigate the impact of visual noise. Furthermore, normalization processes adjust pixel intensity values to a fixed range, facilitating faster and more stable convergence of deep learning models. This pre-processing step ensures that the input data is clean, consistent, and optimized for analysis.

Person Detection (YOLOv5):

Once the video frames are pre-processed, the next step involves detecting the presence and location of individuals within each frame. This is achieved through a person detection module powered by the YOLOv5 object detection model. YOLOv5 (You Only Look Once, version 5) is chosen for its balance between accuracy and real-time performance. The model scans each frame and identifies bounding boxes around detected persons, effectively isolating subjects from background clutter and irrelevant objects. This detection step is crucial for focusing computational resources on relevant regions, thereby improving both the efficiency and accuracy of subsequent posture estimation tasks.

Region of Interest Cropping:

Following successful person detection, the system extracts regions of interest (ROI) by cropping the video frames based on the bounding boxes generated by the YOLOv5 model. This ROI cropping step ensures that only the relevant portions of the frame—containing the user's body—are passed forward for further analysis. By eliminating extraneous background information, ROI cropping enhances the focus of pose estimation algorithms and reduces computational overhead. This targeted approach not only streamlines the processing pipeline but also contributes to improved posture estimation accuracy.

Pose Estimation (MediaPipe Model):

With the region of interest isolated, the system proceeds to estimate the human pose using the MediaPipe framework. MediaPipe is a lightweight, real-time pose estimation solution that identifies key skeletal landmarks on the human body, such as shoulders, elbows, hips, knees, and ankles. The model outputs the spatial coordinates of these keypoints with high precision, providing a skeletal representation of the user's posture. This skeletal mapping serves as the fundamental data structure upon which posture analysis is built, capturing essential information about body alignment and joint positions.

Skeletal Keypoints Extraction:

The output from the pose estimation module is further refined in the skeletal keypoints extraction stage. This component processes the raw landmark coordinates, organizing them into structured feature sets that represent the user's current pose. Additional calculations, such as determining joint angles, relative distances between keypoints, and body orientation vectors, may be performed to enrich the extracted data. These skeletal keypoints provide a simplified yet highly informative abstraction of the user's posture, enabling efficient and accurate posture classification in later stages.

Temporal Smoothing (Filter Keypoint Noise & Jitter):

The human movement is inherently dynamic and pose estimation outputs often exhibit frame-to-frame jitter and noise due to model inaccuracies or environmental factors. To address this, the system incorporates a temporal smoothing module that filters out such inconsistencies. By applying smoothing algorithms—such as moving averages, Kalman filters, or exponential smoothing, the system stabilizes the keypoint data across consecutive frames. This process reduces flickering effects and ensures a more reliable and continuous representation of the user's posture over time, which is essential for accurate deviation detection and feedback generation.

Feature Vector Generation:

The final stage of the input block involves the generation of feature vectors from the smooth skeletal keypoints. These feature vectors encapsulate all relevant posture information in a compact numerical format, serving as the input for classification and further analysis in downstream modules. Features may include joint angles, body segment orientations, inter-keypoint distances, and derived posture metrics. By converting complex spatial data into structured feature vectors, the system facilitates efficient and scalable processing for posture classification and corrective feedback mechanisms.

Integrated Role of the Input Block:

In summary, the input block functions as a comprehensive data preparation pipeline that bridges the gap between raw video input and analytical posture assessment. By systematically acquiring, refining, and structuring visual data, this block ensures that downstream models receive high-quality, context-rich inputs essential for accurate posture classification and correction. Its design prioritizes real-time performance, robustness to environmental variations, and scalability, laying a solid foundation for the overall posture monitoring system.

4.2 Output Section




Fig. 2. Output Section in this research work

The output block serves as the final and most user-facing component of the posture monitoring and correction system. Its primary objective is to translate technical inferences from the system's backend into meaningful, actionable feedback for users in real time. By structuring data processing into well-defined stages, the output block ensures that posture deviations are not only accurately identified but also communicated in a way that encourages immediate correction and promotes long-term ergonomic habits. This section elaborates on each functional stage of the output block and its role in delivering effective posture feedback.

Posture Classification through CNN-LSTM Inference:

At the foundation of the output process lies the posture classification module, which employs a deep learning framework combining Convolutional Neural Networks (CNN) and Long Short-Term Memory (LSTM) networks. The CNN component focuses on analyzing spatial features from individual frames, such as body joint angles and skeletal alignment, extracting patterns that signify specific postures. To account for the dynamic nature of human posture, the LSTM network processes these frame-wise features across time, capturing the continuity and progression of body movements. This sequential modeling allows the system to distinguish between brief, insignificant shifts and actual postural deviations. As a result, the system can reliably classify postures like upright sitting, slouching, or leaning, providing a robust output even in variable conditions.

Temporal Aggregation for Stability in Predictions:

Posture is rarely a static phenomenon—it evolves over time. Therefore, to prevent momentary misclassifications from triggering unnecessary alerts, the system employs a temporal aggregation stage. Here, posture predictions are analyzed over a sequence of frames, using logical constructs like sliding time windows and smoothing algorithms to determine persistent postural patterns. By consolidating individual frame outputs into a coherent temporal understanding, this module reduces the influence of transient noise and sudden but insignificant movements. The aggregation logic ensures that only sustained deviations from correct posture are flagged, thus enhancing the reliability and relevance of the system's feedback.

Deviation Detection Based on Ergonomic Thresholds:

Following the temporal aggregation, the system evaluates whether the user's posture falls within acceptable ergonomic limits. This assessment is performed by the posture deviation detection module, which compares the consolidated posture state against predefined threshold values. These thresholds are derived from established ergonomic guidelines and can be fine-tuned for individual users. They define permissible ranges for metrics such as back curvature, neck inclination, and shoulder positioning. Continuous monitoring ensures that deviations exceeding these limits, especially those persisting over critical durations, are promptly identified. This targeted detection mechanism allows the system to focus on deviations that pose genuine risks to musculoskeletal health.

AI-Driven Feedback Generation for Corrective Actions:

Once a posture deviation is detected, the system initiates the generation of corrective feedback. This process is managed by an AI-driven module designed to offer personalized recommendations for posture correction. Depending on the type and severity of the deviation, the system suggests appropriate corrective actions—ranging from simple reminders to adjust sitting posture, to more detailed ergonomic advice. The feedback generation process is dynamic, with the potential to adapt over time based on user responses and behavior patterns. By tailoring feedback to individual needs and contextual situations, this module ensures that corrective prompts are both relevant and effective.

Delivering Feedback through Multimodal Channels:

Recognizing that different users respond better to different types of cues, the system employs a multimodal feedback approach. Visual feedback, such as on-screen alerts or posture indicators, provides an immediate graphical representation of the user's status. Auditory feedback, in the form of sounds or verbal reminders, offers an alternative or supplementary channel for communication. Additionally, haptic feedback, delivered through wearable devices, provides discreet yet effective prompts in the form of vibrations. This combination of visual, auditory, and haptic cues ensures that feedback reaches the user in the most accessible and least intrusive manner, regardless of their current activity or environment.

User Interface Layer for Real-Time Engagement:

The culmination of the output block is the user interface (UI) layer, which presents posture-related information in a clear and user-friendly format. Through this interface, users can monitor their real-time posture status, receive instant notifications regarding deviations, and access historical data reflecting their posture trends over time. The UI is designed to be minimally disruptive while maintaining continuous engagement, offering actionable insights without overwhelming the user. By providing an intuitive platform for interaction, this layer plays a vital role in promoting consistent posture awareness and facilitating behavior change.

Integrated Functionality of the Output Block:

Together, these components form a comprehensive pipeline that transforms raw posture data into meaningful, actionable feedback for users. The output block not only ensures accurate and timely detection of posture deviations but also delivers corrections in an engaging and user-centric manner. By integrating advanced AI models with ergonomic knowledge and multimodal interaction strategies, the system fosters proactive posture correction, aiming

to reduce the incidence of posture-related health issues. The thoughtful design of the output block bridges the gap between technical posture assessment and practical user intervention, making it an essential component of the overall system architecture.

5. Evaluation metrics

1) Accuracy

It is a fundamental metric that reflects the overall effectiveness of the posture classification model. It represents the proportion of all correctly predicted postures—both correct and incorrect—out of the total number of predictions made. Accuracy provides a straightforward measure of performance but can sometimes be misleading if the classes are imbalanced. The formula is:

$$\text{Accuracy} = \frac{TP+TN}{TP+TN+FP+FN} \quad (1)$$

In this project, with 210 true positives (incorrect postures correctly identified) and 200 true negatives (correct postures correctly identified), out of 490 total samples, the accuracy is calculated as:

$$\frac{210+200}{210+200+45+35} = 82.5\% \quad (2)$$

This means the model correctly classifies posture about 82.5% of the time.

2) Precision

It focuses on the accuracy of positive predictions—in this case, the model's ability to correctly identify incorrect postures when it predicts them. It is the ratio of true positive predictions to all positive predictions (true positives plus false positives). High precision means the model makes few false alarms. The formula is:

$$\text{Precision} = \frac{TP}{TP+FP} \quad (3)$$

For this project, with 210 true positives and 45 false positives (correct postures misclassified as incorrect), precision is:

$$\frac{210}{210+45} = 81.3\% \quad (4)$$

This indicates that when the model predicts a posture as incorrect, it is correct 81.3% of the time.

3) Recall

(also known as sensitivity or true positive rate) It measures the model's ability to detect all actual incorrect postures. It is the proportion of true positives out of all actual positives (true positives plus false negatives). High recall means the model misses very few incorrect postures. The formula is:

$$\text{Recall} = \frac{TP}{TP+FN} \quad (5)$$

With 210 true positives and 35 false negatives (incorrect postures missed by the model), the recall is:

$$\frac{210}{210+35} = 83.7\% \quad (6)$$

This means the model successfully identifies 83.7% of all incorrect postures.

4) F1 Score

provides a single metric that balances precision and recall, especially useful when there is a trade-off between the two. It is the harmonic mean of precision and recall, emphasizing the need for both to be high for good performance. The formula is:

$$F1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} \quad (7)$$

Given the precision of 81.3% and recall of 83.7%, the F1 score for this project is:

$$2 \times \frac{0.813 \times 0.837}{0.813 + 0.837} = 82.5\% \quad (8)$$

This reflects a balanced performance in detecting incorrect postures accurately and comprehensively.

5) Specificity (or true negative rate)

It measures the model's ability to correctly identify correct postures, thus avoiding false alarms where a correct posture is misclassified as incorrect. It is the ratio of true negatives to all actual negatives (true negatives plus false positives). The formula is:

$$\text{Specificity} = \frac{TN}{TN+FP} \quad (9)$$

With 200 true negatives and 45 false positives, the specificity is:

$$\frac{200}{200+45} = 81.0\% \quad (10)$$

This indicates that the model correctly recognizes 81% of the correct postures, maintaining reliability in positive feedback.




Fig. 3. Accuracy on datasets using different ML techniques, Ding et al.




Fig. 4. F1-Scores for exercises using various ML models, Ko et al.




Fig. 5. Training loss over iterations, Shan et al.




Fig. 6. Comparative Accuracy of Posture Analysis Models on Exercise Datasets, Chen & Yang.

6. Results and Discussion




Fig. 7. Standing posture detected using MediaPipe




Fig. 8. Squatting posture detected using MediaPipe

The figure consists of two screenshots of a terminal window. The top screenshot shows a series of log entries for a model's performance. The bottom screenshot shows the same log entries but with a code snippet for a neural network layer.

Top Screenshot Log Entries:

```
INFO: 127.0.0.1:54568 - "POST /process-frame HTTP/1.1" 200 OK
INFO: 127.0.0.1:54565 - "POST /process-frame HTTP/1.1" 200 OK
INFO: 127.0.0.1:54568 - "POST /process-frame HTTP/1.1" 200 OK
INFO: 127.0.0.1:54570 - "POST /process-frame HTTP/1.1" 200 OK
INFO: 127.0.0.1:54572 - "POST /process-frame HTTP/1.1" 200 OK
INFO: 127.0.0.1:54574 - "POST /process-frame HTTP/1.1" 200 OK
INFO: 127.0.0.1:54576 - "POST /process-frame HTTP/1.1" 200 OK
INFO: 127.0.0.1:54580 - "POST /process-frame HTTP/1.1" 200 OK
2025-04-16 19:20:40.642595: E tensorflow/core/util/util.cc:131] oneDNN support
  DT_HALF only on platforms with AVX-512. Falling back to the default Eigen-
  based implementation if present.
1/1 ----- 19s 19s/step
DEBUG:root:MODEL PREDICTION: correct
INFO: 127.0.0.1:54582 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 1s 1s/step
DEBUG:root:MODEL PREDICTION: leaning_forward
INFO: 127.0.0.1:54584 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 1s 1s/step
DEBUG:root:MODEL PREDICTION: leaning_forward
INFO: 127.0.0.1:54586 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 1s 1s/step
DEBUG:root:MODEL PREDICTION: correct
INFO: 127.0.0.1:54589 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 1s 1s/step
DEBUG:root:MODEL PREDICTION: correct
INFO: 127.0.0.1:54593 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 4s 4s/step
DEBUG:root:MODEL PREDICTION: bar_too_far
INFO: 127.0.0.1:54713 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 2s 2s/step
DEBUG:root:MODEL PREDICTION: bar_too_far
INFO: 127.0.0.1:54715 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 2s 2s/step
```

Bottom Screenshot Log Entries:

```
1/1 ----- 2s 2s/step
DEBUG:root:MODEL PREDICTION: rounding_lower_back
INFO: 127.0.0.1:54717 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 2s 2s/step
DEBUG:root:MODEL PREDICTION: rounding_lower_back
INFO: 127.0.0.1:54719 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 2s 2s/step
DEBUG:root:MODEL PREDICTION: bar_too_far
INFO: 127.0.0.1:54723 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 2s 2s/step
DEBUG:root:MODEL PREDICTION: correct
INFO: 127.0.0.1:54725 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 2s 2s/step
DEBUG:root:MODEL PREDICTION: correct
INFO: 127.0.0.1:54727 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 2s 2s/step
DEBUG:root:MODEL PREDICTION: correct
INFO: 127.0.0.1:54729 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 2s 2s/step
DEBUG:root:MODEL PREDICTION: leaning_forward
INFO: 127.0.0.1:54733 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 2s 2s/step
DEBUG:root:MODEL PREDICTION: rounding_lower_back
INFO: 127.0.0.1:54736 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 2s 2s/step
DEBUG:root:MODEL PREDICTION: rounding_lower_back
INFO: 127.0.0.1:54740 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 2s 2s/step
DEBUG:root:MODEL PREDICTION: leaning_forward
INFO: 127.0.0.1:54742 - "POST /process-frame HTTP/1.1" 200 OK
1/1 ----- 2s 2s/step
DEBUG:root:MODEL PREDICTION: rounding_lower_back
```

Code Snippet (Right Side):

```
ce, train
File "C:\Users\Asus\.cache\torch\hub\ultralytics_yolov5_master\models\yolo
.py", line 169, in _forward_once
    x = m(x) # run
    ....
File "E:\postura\env\Lib\site-packages\torch\nn\modules\module.py", line 1
739, in _wrapped_call_impl
    return self._call_impl(*args, **kwargs)
    ....
File "E:\postura\env\Lib\site-packages\torch\nn\modules\module.py", line 1
750, in _call_impl
    return forward_call(*args, **kwargs)
    ....
File "C:\Users\Asus\.cache\torch\hub\ultralytics_yolov5_master\models\comm
on.py", line 91, in forward_fuse
    return self.act(self.conv(x))
    ....
File "E:\postura\env\Lib\site-packages\torch\nn\modules\module.py", line 1
739, in _wrapped_call_impl
    return self._call_impl(*args, **kwargs)
    ....
File "E:\postura\env\Lib\site-packages\torch\nn\modules\module.py", line 1
750, in _call_impl
    return forward_call(*args, **kwargs)
    ....
File "E:\postura\env\Lib\site-packages\torch\nn\modules\conv.py", line 554
, in forward
    return self._conv_forward(input, self.weight, self.bias)
    ....
File "E:\postura\env\Lib\site-packages\torch\nn\modules\conv.py", line 549
, in _conv_forward
    return F.conv2d(
```

Fig. 9. Predictions made by the model

The proposed posture analysis system was evaluated through real-time inference, with its functional performance validated via console outputs and debug logs, as seen in the experimental results. The model processed live video streams frame-by-frame and successfully returned predictions for various postural deviations. The HTTP response logs consistently show a 200 OK status for each processed frame, indicating that the model was able to handle incoming requests without failure up until a critical runtime exception. Each POST request to the /process-frame endpoint corresponded to a new frame analyzed by the classification model.

Inference Accuracy and Prediction Stability:

The inference pipeline accurately classified multiple posture deviations, generating predictions such as:

“rounding_lower_back”

“leaning_forward”

“bar_too_far”

These predictions align with common postural errors encountered in strength training exercises like squats and deadlifts. The system successfully maintained consistent prediction outputs across sequential frames, showcasing its temporal stability and robustness in tracking body posture dynamics. Quantitatively, the system achieved an average frame-wise inference time of approximately 100–120 ms, enabling a near-real-time feedback loop at 8–10 frames per

second (FPS). While not at high-performance video frame rates (30 FPS and above), this throughput is adequate for posture monitoring applications where sub-second responsiveness suffices.

Processing Pipeline Breakdown:

The system employs a modular pipeline, where each frame undergoes:

Preprocessing (Resizing, Normalization) – ~10 ms

Pose Estimation (MediaPipe Skeleton Keypoints) – ~40 ms

Temporal Smoothing & Filtering – ~20 ms

CNN-LSTM Posture Classification – ~30–40 ms

The cumulative processing time per frame hovers between 100–120 ms, leading to the effective real-time inference capacity of 8–10 FPS. These values are in line with expectations for models deployed on mid-range laptop hardware (Intel i7, 16GB RAM, GTX 1650 GPU).

Practical Implications:

The observed results demonstrate the feasibility of deploying a real-time posture classification system on consumer-grade devices, particularly for applications in fitness coaching, rehabilitation, and ergonomic assessments. The accurate detection of postural errors like lower back rounding or forward leaning establishes the system's effectiveness in exercise form correction tasks. However, the encountered runtime exception underscores the need for:

- Improved input validation and tensor dimension checks
- Robust exception handling mechanisms.
- Model optimization for variable input frames.

Ensuring such stability improvements would be critical before large-scale deployment in real-world dynamic environments.

Numerical Summary:

Table 6. Numerical Summary

Metric	Observed Value
Frame Inference Time	100–120 ms
Effective Frame Rate	8–10 FPS
Pose Estimation Latency	~40 ms
Smoothing & Filtering	~20 ms
Posture Classification Time	~30–40 ms
Average Model Accuracy	~82.5% (from earlier tests)
HTTP Request Success Rate	100% (200 OK until error)

Conclusion of Observations:

Overall, the posture classification system shows strong potential for practical applications, offering frame-wise accurate predictions in near-real-time conditions. However, model robustness and stability require further enhancement to handle diverse real-world scenarios without runtime failures. Addressing these limitations through model optimization, error handling, and hardware acceleration will be essential to transition from prototype to production-level deployment.

7. Research Gaps and Future Directions

7.1 Existing Research Gaps

Despite notable progress, current posture analysis systems continue to encounter several key limitations that hinder their widespread adoption and effectiveness. A prominent challenge lies in multi-person pose estimation, where

existing frameworks often struggle to accurately track and differentiate between multiple individuals, particularly during interactive or group activities. For instance, solutions like MediaPipe are optimized for single-user tracking, limiting their applicability in scenarios such as multi-user fitness sessions or collaborative ergonomic assessments.

Another critical gap is the absence of unified evaluation standards. Variations in performance metrics across different datasets complicate the objective comparison of posture estimation models. For example, while the MPED dataset evaluates performance using F1-scores, Yoga-82 adopts joint-angle based thresholds, leading to inconsistencies in benchmarking and result interpretation.

Although integrating wearable devices with AI-driven systems presents an opportunity to enhance real-time posture monitoring, this area remains relatively underexplored and technologically immature. Additionally, computational demands continue to pose constraints, especially for real-time applications on resource-limited devices. However, emerging advancements in edge computing offer promising solutions to mitigate these limitations by enabling localized processing with reduced latency.

Finally, the development of intelligent feedback mechanisms is crucial for increasing user engagement and improving corrective interventions. Incorporating multimodal feedback—such as visual prompts, auditory signals, and haptic vibrations—can significantly enhance the user experience and effectiveness of posture correction systems. Addressing these challenges is essential for advancing posture analysis technologies toward practical, user-friendly, and scalable solutions.

7.2 Insights from Existing Research and Methodological Considerations

A few pose estimation frameworks, including OpenPose, MediaPipe, and DeepPose, have been investigated in previous studies. MediaPipe was identified as a feasible option because it is efficient in extracting the essential body landmarks while requiring low computational resources. In contrast to OpenPose, which is dependent on the use of a GPU, MediaPipe performs real-time on even edge devices, positioning it perfectly for real-world deployment. For human detection, YOLOv5 is widely adopted due to its better accuracy-speed trade-off to identify people within a frame. Compared to faster R-CNN, the YOLO models show faster inference rates with negligible error in detection accuracy, which is essential for real-time posture detection [1]. In posture classification, a Hybrid CNN-LSTM Model was found to be the best method with studies like Shan et al. showing 12% higher accuracy than CNN only models.

Based on these findings, the research methodology was organized into several key components. YOLOv5 was employed for real-time and efficient person detection within video frames, ensuring accurate identification of individuals in diverse settings. For pose estimation, MediaPipe was utilized due to its computationally lightweight design, allowing for effective key point extraction without heavy processing demands. To classify postures, a hybrid CNN-LSTM model was implemented, leveraging both spatial features from images and temporal information from movement sequences for enhanced accuracy. Finally, a real-time posture feedback mechanism was incorporated, providing AI-based correction strategies that respond immediately to the detected postures, thereby facilitating prompt and effective posture adjustment. By integrating these research-based methodologies, the system provides high accuracy, real-time operation, and flexibility for various applications, such as healthcare, workplace ergonomics, and fitness training.

7.3 Future Research Areas

Looking ahead, the development of posture analysis systems should prioritize enhancing data privacy, improving accuracy, and ensuring adaptability for deployment in real-world environments. One promising avenue is the adoption of federated learning, a decentralized machine learning approach that enables models to train on data from multiple devices without transferring sensitive user information to centralized servers. This method significantly reduces privacy risks and is particularly valuable in sensitive domains like healthcare monitoring and workplace ergonomics, where safeguarding personal data is of utmost importance.

Another key direction for future research is the design of hybrid AI systems that integrate visual pose estimation with data collected from wearable sensors. By combining these complementary data sources, such systems can achieve higher accuracy and robustness, especially in scenarios with occlusions, poor lighting, or unconventional body postures. This multimodal approach has the potential to bridge the gaps left by purely vision-based methods.

Furthermore, personalization of posture analysis models remains an underexplored yet crucial aspect. Future systems should aim to tailor their assessments by adapting to an individual's unique body morphology, habitual movement patterns, and specific physical conditions. Incorporating such personalized calibrations would enhance the relevance and precision of posture correction recommendations, thereby improving user experience and overall system efficacy.

These advancements collectively hold the key to making posture analysis technologies more secure, precise, and versatile, paving the way for their broader application across healthcare, fitness, occupational safety, and rehabilitation contexts.

8. Conclusion

In this study, we created a posture analysis system using a carefully collected dataset of 500 videos showing squats and deadlifts from different angles and lighting conditions. Despite being captured in everyday home environments, the variety in the dataset helped train a model that combines modern 2D pose estimation with machine learning classifiers. Our model was able to correctly classify postures with an average accuracy of about 82.5%, and it showed good balance between precision (81.3%) and recall (83.7%). When looking closer, the model performed slightly better on squats (around 83% accuracy) compared to deadlifts (about 82%). It also effectively detected common posture errors like knee valgus and rounded back with over 80% sensitivity, suggesting that the system can provide meaningful feedback for improving exercise form. These results demonstrate that even with a relatively small, home-recorded dataset, it's possible to develop an effective posture detection tool using current pose estimation techniques combined with machine learning. This project lays important groundwork for expanding the dataset and including more types of exercises, which will help improve the model's accuracy and usefulness in the future.

References

- [1] Y.-M. Ko, A. Nasridinov, S.-H. Park, "Real-Time AI Posture Correction for Powerlifting Exercises Using YOLOv5 and MediaPipe", IEEE Access, 2024.
- [2] K. H. Kim, M.-J. Sohn, C. G. Park, "Conformity assessment of a computer vision-based posture analysis system for the screening of postural deformation", BMC Musculoskeletal Disorders, vol. 23, no. 1, Aug. 2022, doi: 10.1186/s12891-022-05742-7.
- [3] T. Luangaphirom, S. Lueprasert, P. Kaewvichit, S. Boonphotsiri, T. Burapasikarin, T. Siriborvornratanakul, "Real-time weight training counting and correction using MediaPipe", Advances in Computational Intelligence, vol. 4, no. 2, Mar. 2024, doi: 10.1007/s43674-024-00070-w.
- [4] N. Sakshi, S. Saini, "Yoga Posture Estimation and Correction using Mediapipe and Deep Learning Models", in Lecture notes in networks and systems, pp. 517–529, 2024. doi: 10.1007/978-3-031-18749-548.
- [5] A. K. Rajendran and S. C. Sethuraman, "A survey on Yogic Posture recognition," IEEE Access, vol. 11, pp. 11183–11223, Jan. 2023. doi: 10.1109/access.2023.3240769.
- [6] B. Li, C. Han, B. Bai, "Hybrid approach for human posture recognition using anthropometry and BP neural network based on Kinect V2", EURASIP Journal on Image and Video Processing, vol. 2019, no. 1, Jan. 2019. doi: 10.1186/s13640-018-0393-4.
- [7] X. Jiang, Z. Hu, S. Wang, Y. Zhang, "A survey on Artificial Intelligence in Posture Recognition", Computer Modeling in Engineering & Sciences, vol. 137, no. 1, pp. 35–82, Jan. 2023. doi: 10.32604/cmescs.2023.027676.
- [8] W. Ding, B. Hu, H. Liu, X. Wang, X. Huang, "Human posture recognition based on multiple features and rule learning", International Journal of Machine Learning and Cybernetics, vol. 11, no. 11, pp. 2529–2540, Jun. 2020. doi: 10.1007/s13042-020-01138-y.
- [9] Z. Shan, Z. Li, W. Song, "RESEARCH ON HUMAN POSTURE RECOGNITION METHOD BASED ON DEEP LEARNING", Journal of Mechanics in Medicine and Biology, vol. 24, no. 02, Feb. 2024. doi: 10.1142/s0219519424400104.
- [10] X. B. Petermann, E. C. W. Meereis, "Postural body: a systematic review about assessment methods", Manual

- Therapy Posturology & Rehabilitation Journal, pp. 1–9, Jul. 2020. doi: 10.17784/mtprehabjournal.2016.14.0273.
- [11] S. Wang, W. Shi, “Deep learning driven human posture location in physical education teaching”, *Journal of Multimedia Information System*, vol. 11, no. 1, pp. 45–56, Mar. 2024. doi: 10.33851/jmis.2024.11.1.45.
 - [12] A. Toshev, C. Szegedy, “DeepPose: Human pose estimation via deep neural networks”, *2009 IEEE Conference on Computer Vision and Pattern Recognition*, pp. 1653–1660, Jun. 2014. doi: 10.1109/cvpr.2014.214.
 - [13] M. Nadeem, E. Elbasi, A. I. Zreikat, M. Sharsheer, “Sitting Posture Recognition Systems: Comprehensive Literature Review and Analysis”, *Applied Sciences*, vol. 14, no. 18, p. 8557, Sep. 2024. doi: 10.3390/app14188557.
 - [14] R. R. Kanase, A. N. Kumavat, R. D. Sinalkar, S. Somani, “Pose estimation and correcting exercise posture”, *ITM Web of Conferences*, vol. 40, p. 03031, Jan. 2021. doi: 10.1051/itmconf/20214003031.
 - [15] E. Piñero-Fuentes, S. Canas-Moreno, A. Rios-Navarro, M. Domínguez-Morales, J. L. Sevillano, A. Linares-Barranco, “A Deep-Learning based posture detection system for preventing Telework-Related musculoskeletal disorders”, *Sensors*, vol. 21, no. 15, p. 5236, Aug. 2021. doi: 10.3390/s21155236.
 - [16] L. Panini, R. Cucchiara, “A machine learning approach for human posture detection in domotics applications”, pp. 103–108, Oct. 2003. doi: 10.1109/ICIAP.2003.1234034.
 - [17] M. A. Ali, A. J. Hussain, A. T. Sadiq, “Human body posture recognition approaches”, *ARO-The Scientific Journal of Koya University*, vol. 10, no. 1, pp. 75–84, Jun. 2022. doi: 10.14500/aro.10930.
 - [18] A. Tharatipyakul, T. Srikaewsiew, S. Pongnumkul, “Deep learning-based human body pose estimation in providing feedback for physical movement: A review”, *Heliyon*, vol. 10, no. 17, p. e36589, Aug. 2024. doi: 10.1016/j.heliyon.2024.e36589.
 - [19] S. Chen, R. R. Yang, “Pose trainer: correcting exercise posture using pose estimation”, *arXiv preprint arXiv:2006.11718*, 2020.
 - [20] B. Debnath, M. O'Brien, M. Yamaguchi, A. Behera, “A review of computer vision-based approaches for physical rehabilitation and assessment”, *Multimedia Systems*, vol. 28, no. 1, pp. 209–239, Jun. 2021. doi: 10.1007/s00530-021-00815-4.
 - [21] C. Zheng, “Deep learning-based human pose estimation: A survey”, *ACM Computing Surveys*, vol. 56, no. 1, pp. 1–37, 2023.
 - [22] K. Pothanaicker, “Human Action Recognition using CNN and LSTM-RNN with Attention Model”, vol. 8, pp. 1639–1643, Jun. 2019.