

Green Agent Analytics: Optimizing Continuous ML Pipelines via Provenance-Based Graphlet Prediction and Cooperative Verification

Sahil Bhoje¹, Vardhak Kore², Purva Vajire³ Dr.A. M. Bagade⁴

¹Student, Pune Institute of Computer Technology, IT, Pune, Maharashtra, India
sahilbhoje15@gmail.com

²Student, Pune Institute of Computer Technology, IT, Pune, Maharashtra, India
vardhakkore@gmail.com

³Student, Pune Institute of Computer Technology, IT, Pune, Maharashtra, India
vajirepurva@gmail.com

⁴ Assistant Professor; SCTR's Pune Institute of Computer Technology, IT, Pune, Maharashtra, India,
ambagade@ms.pict.edu

Abstract

This paper introduces the Resource-Optimized, Trustworthy Multi-Agent System (ROT-MAS), a new framework using Large Language Models (LLMs) to fully automate the data science pipeline. While current systems exist, they are not reliable enough for industrial use because they are often inefficient and do not keep a clear audit trail. This study solves these problems with a three-part approach. First, a Multi-Agent Debate (MAD) allows agents to critique each other's plans to find and fix errors before they run. Second, an LLM-guided Provenance Tracker creates a clear audit trail and offers multiple "granularity" levels (like 'Full' or 'Only Columns') to save storage space. Third, a Wasted Computation Mitigation (WCM) module predicts and stops unnecessary model training, making the process "Green." The results indicate that the ROT-MAS is very effective: the MAD protocol reduced critical errors by 72%, the WCM module cut computational waste by 53%, and the "Only Columns" provenance tracker reduced storage by 99.4% while the automated documentation achieved 98% precision. Keywords: Data Provenance, Automated Data Science, Multi-Agent Systems, Green AutoML, LLM Agents, Multi-Agent Debate (MAD), and Wasted Computation Mitigation (WCM).

Keywords: Graphlet prediction, Multi-agent debate, Large language model, Provenance tracker, wasted computation mitigation

1. Introduction

Integrating Large Language Models (LLMs) into intelligent agents is a key step toward AGI [2], [4]. These agents are starting to automate complex data science tasks [10], [20], which helps make machine learning accessible to non-experts [22], [24]. However, current systems are not ready for production use. They suffer from three critical problems: (1) Safety Risks, as LLMs can hallucinate and are insecure [2], [14]; (2) No Audit Trail, making it difficult to track their decisions (provenance) for compliance [25], [26]; and (3) High Computational Waste, as they are slow and many trained models are never deployed [13], [27].

To solve these issues, this paper introduces the Resource-Optimized, Trustworthy Multi-Agent System (ROT-MAS). The proposed is the first to combine three specific solutions to these problems: a debate-based mechanism for reliability, a scalable provenance tracker for auditing, and a two-tier optimization strategy to reduce cost and waste.

2. Related Work

2.1. Fundamental Ideas in Agent Design

Current research is focused on building reliable LLM agents. This "Agent Methodology" has two main components: memory and planning. To be autonomous, agents need long-term memory (LTM) to store and recall past experiences, such as in the tiered architecture of Mem GPT [6]. They must also learn from their actions using methods like self-reflection [7].

Agent planning has also become more advanced. It has moved from simple, linear steps to complex, multi-step reasoning. For example, Tree-of-Thought (Tot) [13] allows an agent to explore different reasoning paths in a tree structure and backtrack if needed. Graph-of-Thought (Got) [1] makes this concept even more flexible by using a graph structure.

To improve reliability, the field is moving from single-agent to multi-agent systems. Frameworks like Multi-Agent Debate (MAD) use cooperation and disagreement between agents to find errors and reduce hallucinations. This is part of a larger trend of using ML engineering principles to make agents transparent and dependable. This includes data provenance for auditing, security guardrails, and efficiency optimizations like LORA. New tools like Agent Bench are also being developed to properly evaluate these new, complex agent systems.

Table 1. Comparative Literature Survey

Sr. No	Author(s)	Journal/Conference	Key Findings	Gaps
1.	Besta et al. (2024) [1]	AAAI Conference	Introduces Graph of Thought (GoT), a flexible planning architecture that models thoughts as a graph ¹⁵¹⁵¹⁵ .	Primarily a planning model. Lacks mechanisms for external critique (MAD) or predictive resource optimization (WCM).
2.	Huang et al. (2024) [2]	arXiv	A comprehensive survey on LLM agent planning, covering memory, reflection, and guardrails ¹⁶ .	A survey paper; does not propose a specific, integrated, and auditable system for data science pipelines.
3.	Long (2023) [3]	arXiv	Proposes a "Large language model guided tree-of-thought" ¹⁷¹⁷¹⁷ .	Focuses on single-agent reasoning paths. Does not incorporate multi-agent debate (MAD) for error correction.
4.	Luo et al. (2025) [4]	arXiv	A survey on LLM Agent Methodology, Applications, and Challenges ¹⁸ .	A survey paper; identifies challenges but does not provide an operational, resource-optimized solution.
5.	Ma et al. (2024) [5]	arXiv	A survey on vision-language-action models for embodied AI ¹⁹ .	Focused on embodied AI. Not applicable to the <i>data science pipeline</i> domain, which requires data provenance.
6.	Packer et al. (2023) [6]	arXiv	Introduces MemGPT, an agent architecture with tiered memory for long-term context and self-reflection ²⁰²⁰²⁰ .	Solves memory, but not pre-execution reliability (MAD) or auditable data lineage (Provenance).
7.	Shinn et al. (2023) [7]	arXiv	Introduces Reflexion, an agent that uses dynamic memory and self-reflection to learn from past	Self-reflection is post-hoc. Lacks a <i>proactive</i> multi-agent debate (MAD) to prevent errors <i>before</i>

			actions ²¹²¹²¹²¹ .	execution.
8.	Shinn et al. (2023) [8]	NeurIPS	Formalizes Reflexion with verbal reinforcement learning ²² .	Same as [7]; focused on post-hoc learning, not proactive verification or resource efficiency (WCM).
9.	Sun et al. (2023) [9]	NeurIPS	Proposes AdaPlanner, an adaptive planner that learns from feedback ²³ .	Focuses on adapting the <i>plan</i> . Does not track the <i>data's</i> lineage (Provenance) or predict computational waste (WCM).
10.	Wang et al. (2024) [10]	TMLR	Introduces Voyager, an open-ended embodied agent that builds a skill library from experience ²⁴²⁴²⁴ .	Domain-specific (Minecraft) ²⁵ . Lacks the specific data science auditability (Provenance) and safety (MAD) required for production.
11.	Wang et al. (2023) [11]	arXiv	Proposes interactive planning for open-world multi-task agents ²⁶ .	Relies on human-in-the-loop for planning. Does not have an <i>automated</i> multi-agent debate (MAD) or predictive efficiency (WCM).
12.	Wang et al. (2024) [12]	NAACL	Proposes RecMind, an LLM agent specifically for recommendation systems ²⁷ .	Domain-specific. Does not address generalized data science pipelines, nor does it include a Green AutoML (WCM) component.
13.	Yao et al. (2023) [13]	NeurIPS	Introduces Tree of Thoughts (ToT), a planning method for deliberate problem-solving using a tree structure ²⁸²⁸²⁸²⁸ .	A foundational single-agent planning method. Does not include multi-agent critique (MAD) or data auditability (Provenance).
14.	Zhao et al. (2024) [14]	AAAI Conference	Introduces Expel, an LLM agent that learns from experiences ²⁹ .	Focuses on learning from <i>past</i> experience. Does not have a <i>real-time</i> debate (MAD) or a <i>predictive</i> efficiency module (WCM).
15.	Zhu et al. (2023) [15]	arXiv	Proposes GHOST, a generally capable agent for open-world environments (Minecraft) ³⁰ .	Domain-specific. Lacks the production-grade features of auditable provenance and computational waste reduction (WCM).
16.	Nejjar et al. (2023) [16]	arXiv	Explores the use of LLMs for code generation and data analysis in scientific contexts ³¹ .	Focuses on code generation. Does not propose a holistic, <i>trustworthy</i> (MAD) or <i>auditable</i> (Provenance) agent system.
17.	Hong et al. (2024) [17]	arXiv	Introduces Data Interpreter, an LLM agent specifically for data science tasks ³² .	An end-to-end system, but lacks the specific optimization modules of ROT-MAS (MAD for reliability, WCM for efficiency).

18.	Yang et al. (2024) [18]	arXiv	Explores LLMs for predictive tabular tasks in data science ³³ .	Focuses on the <i>model training</i> aspect. Does not propose a full <i>agent system</i> to manage the pipeline, nor does it address auditability.
19.	Nagarajah & Poravi (2019) [19]	IEEE 5th I2CT	A review on Automated Machine Learning (AutoML) systems ³⁴ .	A review of traditional AutoML. Does not cover LLM-driven <i>agents</i> or address their specific gaps (hallucination, provenance).
20.	Chugh et al. (2023) [20]	Vol. 30, pp. 152-157	Discusses "Intelligent agents driven data analytics using large language models" ³⁵ .	A conceptual paper; lacks the specific, implemented modules for trust (MAD), auditability (Provenance), and efficiency (WCM).

3. Proposed Methods

Here, Architecture is described the optimized employing sophisticated agent coordination and data management frameworks, emphasizing dependability, transparency, and computational economy. EDA, preprocessing, and autoML basic components were used by the system, but they were integrated into a self-optimizing ecosystem.

3.1 Overall System Architecture

The system flow diagram is shown in Fig. 1. It is explained stepwise as follows:

Step 1: The user submits a "User Request."

Step 2: The request goes to a "Planner Agent," which creates an "Initial Plan."

Step 3: The "Multi-Agent Debate (MAD)" module reviews and refines this plan into a "Verified Plan."

Step 4: The "Predictive WCM" module evaluates the "Verified Plan" and makes a "Go/No-Go Decision."

Step 5: Following a "Go" decision, the "Executor Agent" begins to carry out the sub-tasks.

Step 6: The "Executor Agent" uses "External Tools" to help complete the tasks.

Step 7: As the system works, the "Multi-Granular Provenance Tracker" records all actions.

Step 8: Finally, the system produces the "Final Output," such as a report or model.

3.2 Fundamental Agent Functions (Hierarchical Model)

The system utilized a hybrid, hierarchical architecture that assigned execution to specialized agents through a Centralized Coordinator. The Interactive AI Chatbot served as the User Proxy and main interface, acting as the Human Feedback Input (HFI) mechanism to translate user natural language requests into organized tasks. This workflow was overseen by the Central Coordinator Agent, which maintained the global state and tracked the specialized agents' confidence scores, where any low-confidence assignment triggered the Verification Layer. Additionally, Executive Agents responsible for EDA, Preprocessing, and Training utilized PyCaret, pandas, and

scikit-learn to execute specific tasks, ensuring that all data transformation scripts were communicated to the Provenance Tracker Agent by the Preprocessing Agent prior to execution.




Fig. 1. ROT-MAS System Workflow Diagram

3.3 The MAD Protocol (Multi-Agent Debate)

A mandatory multi-agent discussion mechanism was implemented for validation, which took place prior to the execution of crucial stages, in order to maximize reliability and combat LLM hallucinations.

3.3.1 Overview of Concepts

A three-agent debate protocol was initiated when the Coordinator identified a high-risk operational plan.

1. **Executive Agent:** Outlined the suggested course of action or script.
2. **Critic Agent:** Serving as an antagonistic check, the agent evaluated the plan for possible defects, prejudices, or ethical issues.
3. **Reflector Agent:** Performed self-correction by going over the discussion, accessing internal memory (LTM) of comparable previous actions, and offering a final, improved recommendation. The outcome was sent to the user through the Chatbot for Human-Agent Supervision (HAC) if the internal debate did not meet a confidence threshold.

3.3.2 The Pseudocode MAD Algorithm

The following is the MAD protocol's algorithm:

1. Begin
2. Input: is_critical_task (boolean), initial_plan
3. Coordinator.measure_plan_confidence(initial_plan) = confidence_score
4. IF the task is not critical and the confidence score is more than the high threshold Next,
5. RETURN initial_plan // No discussion is required
6. END IF
7. // Start the Debate Protocol
8. ExecutiveAgent.present(initial_plan) = plan_proposal
9. critique = CriticAgent.evaluate(plan_proposal)
10. LTM.consult(plan_proposal.topic) = past_executions
11. refined_plan = ReflectorAgent.review(memory=past_executions, debate=[plan_proposal, critique])
12. Coordinator.measure_plan_confidence(refined_plan) = final_confidence
13. If is a significant task or ultimate confidence is less than the confidence threshold Next,
14. HAC Fallback: RETURN HumanAgent.supervise(refined_plan)
15. ELSE
16. RETURN refined_plan // Verified Plan approved
17. END IF
18. END

3.4 The Provenance Tracker with Multiple Groups

To guarantee transparent, auditable data management for data wrangling activities, PROLIT principles were integrated into the system. The Python script of the Preprocessing Agent was automatically analyzed and rewritten by the Provenance Tracker Agent using an LLM, a process that accomplished standardized structure segmentation, the insertion of thorough comments detailing every operation, and the identification of columns used in each activity. Additionally, the tracker produced provenance graph fragments that were stored in a graph database while monitoring execution. To maximize storage efficiency, the capture functioned at four adjustable granularity levels:

- Full Level: All activity, column, and entity details were recorded.
- Sketch Level: A minimum representation that maximizes query performance.
- Derivation Level: To trace data lineage, "was Derived From" linkages were kept.
- Only Columns Level: For a rough overview, only activity and column nodes were included; entities were left out

3.5 The Module for Wasted Computation Mitigation (WCM)

For computational efficiency, a two-level optimization technique was utilized. First, the Wasted Computation Mitigation (WCM) approach was implemented based on Model Graphlet analysis, extracting features such as pipeline shape, model type, and input data similarity from the current graphlet run. Second, using a mechanism for predictive pruning, the WCM Module applied these features to a supervised machine learning model, such as Random Forest, prior to the Training Agent's execution. This allowed the system to pre-emptively skip or down-

prioritize training sessions to avoid wasting compute if the WCM projected a low possibility of deployment, effectively identifying "unpushed graphlets".

3.6 Analysis of Algorithmic Complexity

Resource optimization is a fundamental feature of the ROT-MAS framework, made possible by the low computing complexity of the innovative modules. The time complexity of the MAD protocol is $O(1)$ relative to the dataset size, as the discussion consists of a predetermined number of LLM calls and LTM lookups. Since the complexity depends on the plan metadata rather than the actual data, it is unaffected by the number of rows (N) or features (C), resulting in a space complexity of $O(1)$. Conversely, the complexity of the provenance tracker increases with pipeline steps (P), columns (C), and rows (N):

- Full Level: The space complexity is $O(P * C * N)$ due to the need to store a node or entity for each data point at each stage, resulting in significant storage overhead.
- Only Columns Level: The space complexity is $O(P * C)$. Storage requirements become independent of the dataset length by eliminating row-level entities (N), scaling only with pipeline stages and columns, which enabled a 99.4% reduction in storage.

4. Methodology

To empirically validate the performance of the suggested ROT-MAS framework, a series of tests were conducted to assess the system's dependability, effectiveness, and auditability. In the experimental design, the three main optimization modules were compared against baseline models to measure specific improvements. For reliability, the planning component was evaluated against a typical Chain-of-Thought (Cot) baseline to gauge the decrease in planning failures. For efficiency, the system's predictive pruning was tested by modelling a continuous deployment pipeline to determine the amount of wasted computation, or "unpushed graphlets," that was effectively avoided. Finally, for auditability, the lineage of a standardized, intricate data-wrangling pipeline was recorded at each granularity level to determine the storage footprint of the Multi-Granular Provenance Tracker.

5. Metrics for Performance Evaluation

5.1 Experimental Environment & Benchmark Datasets

Every experiment was carried out on a workstation that had an NVIDIA RTX 3050 GPU, 8GB of RAM, and an Intel Core i5 CPU. The core planner agent used the Llama 3 8B model, while the software environment included Python 3.10, PyCaret 3.1, pandas 2.1, and Neo4j 5.

The agent was required to examine seven common, publicly accessible benchmark datasets in order to verify the effectiveness of the ROT-MAS framework.

5.2 Benchmark Dataset Characteristics: Table 2. Dataset Characteristics

Dataset Name	Domain	No. of Rows	No. of Columns	Task Type
UCI Iris	Biology	150	4	Classification
UCI Wine	Chemistry	178	13	Classification
UCI Breast Cancer	Medical	569	30	Classification
Kaggle Titanic	Tabular	891	12	Classification
UCI Car Evaluation	Tabular	1,728	6	Classification
UCI Abalone	Biology	4,177	8	Regression
Boston Housing	Tabular	506	14	Regression

Critical Error Rate (for MAD):

This statistic measures the frequency of hallucinations or critical defects in a generated plan and is used to assess the 72% decrease in planning errors. It is described as:

$$\text{Critical Error Rate} = \frac{\text{Number of Hallucinated or Flawed Plans}}{\text{Total Plans Generated}} \quad (1)$$

Wasted Computation (for WCM):

This metric determines the percentage of compute time saved by the WCM's predictive pruning and is used to gauge the 53% decrease in computational overhead.

$$\text{Wasted Computation Saved} = \frac{\sum \text{Runtime of Pruned Runs}}{\sum \text{Runtime of All Runs}} \quad (2)$$

Storage Overhead (for Provenance):

The total file size (in MB) needed to store the entire provenance graph for a particular pipeline is used to calculate the 99.4% storage reduction.

$$\text{Storage Overhead} = \text{File Size of Graph Database (MB)} \quad (3)$$

Documentation Fidelity (for Provenance):

Used to gauge the LLM-guided rewriting's 98% accuracy and 95% recall. TP (True Positive) is a properly documented data transformation, according to the standard definitions that is employed.

$$\text{Precision} = \frac{TP}{TP + FP} \quad (4)$$

$$\text{Recall} = \frac{TP}{TP + FN} \quad (5)$$

6. Results & Discussion

The technique outlined in Section IV was used to assess the ROT-MAS framework and its baseline counterparts. The notable improvements in computing efficiency, auditability, and dependability are shown in the following results.

6.1 Reliability Outcomes (MAD Protocol)

The Multi-Agent Debate (MAD) protocol greatly enhanced planning consistency and factual accuracy. When compared directly to a typical Chain-of-Thought (Cot) baseline, the MAD methodology resulted in a 72% reduction in critical errors (hallucinations and false positives) during the data preparation and validation stages.

As seen in Fig. 2, the baseline Cot agent produced a significant rate of critical errors when tasked with evaluating the benchmark datasets. However, most of these errors were detected and self-corrected by the MAD protocol-equipped ROT-MAS before they were executed.

In Fig. 2, the comparison of Critical Error Rate (Baseline vs. MAD), the critical errors of the ROT-MAS (MAD) system (Series 2) and the Baseline (Cot) system (Series 1) are contrasted in this graph over four pipelines. The ROT-MAS system continuously lowers these mistakes, exhibiting an average reduction of 72%, whereas the baseline has a significant error rate.




Fig. 2. Comparison of Critical Error Rate (Baseline vs. MAD)

6.2 Computational Efficiency Results (WCM & LLM)

Two main optimizations produced notable increases in computational efficiency. First, non-impactful model retraining runs were successfully found and pre-emptively pruned by the Predictive Wasted Computation Module (WCM). This change removed 53% of unnecessary computational overhead throughout the ML pipelines under observation, as seen in Fig. 3.

In Figure 3, the Wasted Computation Overhead (Baseline vs. WCM) is represented. The graph contrasts ROT-MAS system (orange bars) with the Baseline system (blue bars) for four distinct data science pipelines. The quantity of computation that is wasted is represented by the height of the bars. The orange bar (ROT-MAS) is consistently significantly shorter than the blue bar (Baseline). These results average out to the 53% total reduction in wasted computation that you claimed, demonstrating the constant effectiveness of your WCM module. Second, inference-level optimizations were very successful. The Interactive AI Chatbot's average inference latency was lowered by 40% by applying the LoRA (Low-Rank Adaptation) fine-tuning technique to the core LLM, greatly enhancing the system's responsiveness.




Fig. 2. Wasted Computation Overhead (Baseline vs. WCM)

In Figure 4, the comparison of LLM Inference Latency (Baseline vs. LoRA-Optimized) is shown. For five different test queries, this graph compares the LoRA-Optimized LLM latency (orange bars) with the Baseline LLM latency (blue bars). In every query, the blue bar (Baseline) is significantly longer than the orange bar (LoRA-Optimized). According to this data, the optimization was successful in cutting the average inference latency by 40%, greatly enhancing the responsiveness of the system.



Fig. 3. LLM Inference Latency (Baseline vs LoRA-Optimized)

6.3 Auditability & Storage Results (Provenance)

Both auditability and storage efficiency were effectively achieved by the Multi-Granular Provenance Tracker. When compared to the thorough "Full Level" tracking, the "Only Columns" tracking level showed impressive optimization, cutting the storage overhead by 99.4%, as illustrated in Fig. 5. A Comparison of Storage Overhead (Full Level vs. Only Columns Level) is presented in Fig. 5. This graph compares the storage costs of your optimized "Only Columns Level" (orange bars) with the "Full Level" provenance (blue bars). The test was carried out using four different data science workflows. The "Full Level" technique uses a lot of storage, as evidenced by the extremely tall blue bar in every test. By contrast, your "Only Columns Level" solution uses extremely little storage because the orange bar is almost undetectable due to its small length. This figure, which consistently achieves the 99.4% storage reduction you mentioned, demonstrates how effective your optimization is.




Fig. 4. Storage Overhead (Full Level vs. Only Columns Level)

Regarding documentation, high fidelity was attained in the automated tasks of the LLM-powered PROLIT component. The automatically generated records and descriptions demonstrated 98% precision and 95% recall when compared to human created documentation as shown in Fig. 6. Here, documentation fidelity (Precision & Recall) This graph shows how well your automated documentation system performs on four different test sequences. The blue line represents precision, or how exact the documentation is, while the orange line represents recall, or how thorough the documentation is. Both lines stay quite high on the chart, at or above 94%, for each test. This shows that your system is incredibly thorough and precise across all of the different processes. The smooth and steady lines, which average out to the 98% Precision and 95% Recall you claimed, show how reliable and stable your system is.




Fig. 5. Documentation Fidelity (Precision and Recall)

The experimental results demonstrate that the ROT-MAS is more than just an automation tool. It is a practical framework for making LLM agents both trustworthy and efficient. It was demonstrated that trust can be "built-in" rather than "bolted-on." The MAD protocol provides internal trust by stopping errors before they happen, which is better than reacting to them later. The Provenance tracker provides external accountability, giving users a complete audit trail without the 99.4% storage cost of full tracking. This framework also creates a "Green Auto ML" system. By using the WCM module to predict and stop unnecessary training runs, the system directly saves significant computation time and money. This is a practical step toward making large-scale AI sustainable.

7. Conclusion

This research successfully built the Resource-Optimized, Trustworthy Multi-Agent System (ROT-MAS), proving that LLM-driven data science can be made production-ready. This research directly addressed major industry gaps by combining three features: (1) a Multi-Agent Debate (MAD) to improve reliability, (2) an LLM-Guided Provenance Tracker for clear auditing, and (3) a Green Auto ML (WCM) strategy to maximize efficiency. This combined ROT-MAS framework provides a crucial blueprint for building future LLM agents that are not just autonomous, but also efficient, reliable, and accountable enough for high-stakes tasks.

7.1 Future Scope

Future work will focus on improving the Human-Agent Collaboration (HAC) component. While the current system allows for human intervention, it is necessary to measure exactly how much effort this requires from non-

technical users using standardized tests. The objective is to shift away from purely autonomous agents. Instead, the aim to build systems where LLMs act as supportive partners that reduce the user's cognitive load. This approach will make AI tools more practical and effective for real-world applications.

References

- [1] M. Besta, N. Blach, A. Kubicek, R. Gerstenberger, L. Gianinazzi, J. Gajda, T. Lehmann, M. Podstawski, H. Niewiadomski, P. Nyczyk, and T. Hoefler, "Graph of thoughts: Solving elaborate problems with large language models," in *Proc. AAAI Conference on Artificial Intelligence*, vol. 38, no. 16, pp. 17682–17690, 2024.
- [2] X. Huang, W. Liu, X. Chen, X. Wang, H. Wang, D. Lian, Y. Wang, R. Tang, and E. Chen, "Understanding the planning of llm agents: A survey," *arXiv preprint arXiv:2402.02716*, 2024.
- [3] J. Long, "Large language model guided tree-of-thought," *arXiv preprint arXiv:2305.08291*, 2023.
- [4] J. Luo, W. Zhang, Y. Yuan, Y. Zhao, J. Yang, Y. Gu, B. Wu, B. Chen, Z. Qiao, Q. Long, R. Tu, X. Luo, W. Ju, Z. Xiao, Y. Wang, M. Xiao, C. Liu, J. Yuan, S. Zhang, Y. Jin, F. Zhang, X. Wu, H. Zhao, D. Tao, P. S. Yu, and M. Zhang, "Large Language Model Agent: A Survey on Methodology, Applications and Challenges," *arXiv preprint arXiv:2503.21460v1*, 2025.
- [5] Y. Ma, Z. Song, Y. Zhuang, J. Hao, and I. King, "A survey on vision-language-action models for embodied ai," *arXiv preprint arXiv:2405.14093*, 2024.
- [6] C. Packer, V. Fang, S. G. Patil, K. Lin, S. Wooders, and J. E. Gonzalez, "Memgpt: Towards llms as operating systems," *arXiv preprint arXiv:2310.08560*, 2023.
- [7] N. Shinn, B. Labash, and A. Gopinath, "Reflexion: an autonomous agent with dynamic memory and self-reflection," *arXiv preprint arXiv:2303.11366*, 2023.
- [8] N. Shinn, F. Cassano, A. Gopinath, K. Narasimhan, and S. Yao, "Reflexion: Language agents with verbal reinforcement learning," in *Advances in Neural Information Processing Systems*, vol. 36, pp. 8634–8652, 2023.
- [9] H. Sun, Y. Zhuang, L. Kong, B. Dai, and C. Zhang, "Adaplaner: Adaptive planning from feedback with language models," in *Advances in Neural Information Processing Systems*, vol. 36, pp. 58202–58245, 2023.
- [10] G. Wang, Y. Xie, Y. Jiang, A. Mandlekar, C. Xiao, Y. Zhu, L. Fan, and A. Anandkumar, "Voyager: an open-ended embodied agent with large language models," *Transactions on Machine Learning Research*, 2024.
- [11] Z. Wang, S. Cai, A. Liu, X. Ma, and Y. Liang, "Describe, explain, plan and select: Interactive planning with large language models enables open-world multi-task agents," *arXiv preprint arXiv:2302.01560*, 2023.
- [12] Y. Wang, Z. Jiang, Z. Chen, F. Yang, Y. Zhou, E. Cho, X. Fan, X. Huang, Y. Lu, and Y. Yang, "Recmind: Large language model powered agent for recommendation," in *Findings of the Association for Computational Linguistics: NAACL 2024*, pp. 4351–4364, 2024.
- [13] S. Yao, D. Yu, J. Zhao, I. Shafra, T. L. Griffiths, and K. Narasimhan, "Tree of thoughts: Deliberate problem solving with large language models," in *Advances in Neural Information Processing Systems*, vol. 36, pp. 11809–11822, 2024.
- [14] A. Zhao, D. Huang, Q. Xu, M. Lin, Y. J. Liu, and G. Huang, "Expel: Llm agents are experiential learners," in *Proc. AAAI Conference on Artificial Intelligence*, vol. 38, pp. 19632–19642, 2024.
- [15] X. Zhu, Y. Chen, H. Tian, C. Tao, W. Su, C. Yang, G. Huang, B. Li, L. Lu, X. Wang, Y. Qiao, Z. Zhang, and J. Dai, "Ghost in the minecraft: generally capable agents for open-world environments via large language models with text-based knowledge and memory," *arXiv preprint arXiv:2305.17144*, 2023.
- [16] M. Nejjar, L. Zacharias, F. Stiehle, and I. Weber, "LLMs for science: Usage for code generation and data analysis," *arXiv preprint*, 2023.
- [17] S. Hong, Y. Lin, B. Liu, B. Liu, B. Wu, C. Zhang, D. Li, J. Chen, J. Zhang, J. Wang, L. Zhang, L. Zhang, M. Yang, M. Zhuge, T. Guo, T. Zhou, W. Tao, R. Tang, X. Lu, X. Zheng, X. Liang, Y. Fei, Y. Cheng, Y. Ni, Z. Gou, and Z. Xu, "Data interpreter: An LLM agent for data science," *arXiv preprint*, Mar. 2024.

- [18] Y. Yang, Y. Wang, S. Sen, L. Li, and Q. Liu, "Unleashing the potential of large language models for predictive tabular tasks in data science," arXiv preprint, Apr. 2024.
- [19] T. Nagarajah and G. Poravi, "A review on automated machine learning (AutoML) systems," in Proc. 2019 IEEE 5th International Conference for Convergence in Technology (I2CT), Mar. 2019.
- [20] T. Chugh, K. Tyagi, R. Seth, and P. Srinivasan, "Intelligent agents driven data analytics using large language models," vol. 30, pp. 152–157, Nov. 2023.
- [21] J. Xu, J. Li, Z. Liu, N. Suryanarayanan, G. Zhou, J. Guo, H. Iba, and K. Tei, "Large language models synergize with automated machine learning," arXiv preprint, 2024.
- [22] P. Trirat, W. Jeong, and S. J. Hwang, "AutoML-agent: A multi-agent LLM framework for full-pipeline AutoML," arXiv preprint, 2024.
- [23] Z. Rasheed, M. Waseem, A. Ahmad, K. K. Kemell, W. Xiaofeng, A. N. Duc, and P. Abrahamsson, "Can large language models serve as data analysts? A multi-agent assisted approach for qualitative data analysis," arXiv preprint, Feb. 2024.
- [24] Z. Shang, E. Zraggen, B. Buratti, F. Kossmann, P. Eichmann, Y. Chung, C. Binnig, E. Upfal, and T. Kraska, "Democratizing data science through interactive curation of ML pipelines," in Proc. 2019 International Conference on Management of Data, Jun. 2019.
- [25] R. J. Urbanowicz, R. Zhang, Y. Cui, and P. Suri, "STREAMLINE: A simple, transparent, end-to-end automated machine learning pipeline facilitating data analysis and algorithm comparison," arXiv preprint, 2022.
- [26] X. Li and T. Döhmen, "Towards efficient data wrangling with LLMs using code generation," arXiv preprint, May 2024.
- [27] A. Moharil, J. Vanschoren, P. Singh, and D. Tamburri, "Towards efficient AutoML: A pipeline synthesis approach leveraging pre-trained transformers for multimodal data," Machine Learning, vol. 113, no. 9, pp. 7011–7053, Jul. 2024.
- [28] R. Barbudo, A. Ramírez, and J. R. Romero, "Evolving machine learning workflows through interactive AutoML," arXiv preprint, 2024.
- [29] Q. Yang, S. Nikolenko, M. Ongpin, I. Gossoudarev, Y. Y. Chu-Farseeva, and A. Farseev, "SOMONITOR: Explainable marketing data processing and analysis with large language models," arXiv preprint, 2024.
- [30] A. Tornede, D. Deng, T. Eimer, J. Giovanelli, A. Mohan, T. Ruhkopf, S. Segel, D. Theodorakopoulos, T. Tornede, H. Wachsmuth, and M. Lindauer, "AutoML in the age of large language models: Current challenges, future opportunities and risks," arXiv preprint, Feb. 2024.